

LINUX FOUNDATIONS

Linux Foundations

Operate the system. Trust the evidence.

Terminal · permissions · SSH · logs · systemd

6-hour workshop

50% theory + demos · 50% hands-on labs

Ubuntu 24.04 browser labs · macOS + Windows



By the end, you can recover your bearings

01 **Navigate**

Move through Linux, manipulate files, and use pipes without losing track of location.

02 **Control access**

Reason about users, groups, ownership, modes, sudo, and SSH keys.

03 **Diagnose**

Start and stop services, filter logs, identify a cause, and verify the repair.

Your browser is a window into another Linux system

BROWSER

**Safari, Edge, or Chrome
on macOS or Windows**

Shows instructions + terminal

`http://linux-lab.local`

SESSION

**Ubuntu 24.04
in an Incus container**

Runs `ls`, `sudo`, `apt`, `systemctl`

`learner@linux-lab:~$`



Browser outside. Linux inside.

Practice gets exactly half of the workshop

180 min

Theory ♦ live demos

- Linux, filesystems, terminal
- Identity, permissions, editor
- SSH, logs, systemd

180 min

Six participant labs

- Open an isolated lab
- Reset, solve, verify state
- Repair a failed service

Breaks can sit around the blocks; the instructional split stays 50 / 50.

Linux is the kernel; a usable system is a stack

Applications

ssh · apt · nano · your services

User space

shells · libraries · GNU tools · systemd

Linux kernel

processes · memory · networking · devices

Hardware / VM / container

CPU · RAM · disks · virtual devices

The kernel turns hardware into abstractions

Processes scheduled programs with isolated address spaces

Memory virtual memory, pages, caches, and protection

Filesystems one namespace across disks and virtual data

Network sockets, routes, interfaces, and packet flow

Devices drivers expose hardware through common interfaces

Kernel space and user space fail differently

KERNEL SPACE

shared authority

- Drivers and core resource control
- A crash can affect the whole system
- Observed with dmesg and kernel journal fields

USER SPACE

isolated processes

- Shells, tools, daemons, applications
- A crash usually affects one process
- Observed with exit status and unit logs

Locate the failing layer before choosing a tool.

A process is a running program with context

PID

a temporary process identifier

Parent

the process that created it

Identity

effective user and groups used for access checks

Environment

name/value data inherited at process start

Descriptors

open files, sockets, and streams

A distribution is a curated operating agreement

Same fundamentals

- Kernel interface
- Filesystem concepts
- Users and permissions
- OpenSSH + systemd

Different choices

- Package format
- Release cadence
- Default configuration
- Support lifecycle

Choose for context

- Ubuntu / Debian
- RHEL / Rocky / Alma
- Fedora
- Cloud/vendor images

Today: Ubuntu 24.04 LTS keeps every lab consistent.

Distribution families share ideas, not every command

Debian family

- Ubuntu · Debian
- APT + .deb
- /etc/apt
- Long support

Red Hat family

- RHEL · Rocky · Alma
- DNF + RPM
- /etc/yum.repos.d
- SELinux defaults

Other contexts

- Fedora: newer stacks
- Alpine: small images
- Arch: rolling
- Vendor clouds

Identify the family before copying package or service instructions.

Stable and rolling releases trade currency

STABLE / LTS

predictability

- Longer support window
- Security fixes with fewer surprises
- Older major versions are common

ROLLING

currency

- Frequent package updates
- Fast access to new features
- More integration change over time

Choose for workload and support needs, not fashion.

Record six facts before following external instructions

Distribution

```
cat /etc/os-release
```

Kernel

```
uname -r
```

Architecture

```
uname -m
```

Package tool

```
command -v apt dnf apk
```

Init system

```
ps -p 1 -o comm=
```

Privilege

```
id; sudo -l
```

Virtual machines isolate kernels; containers share one

VIRTUAL MACHINE

hardware boundary

- Guest owns a kernel
- Different isolation boundary
- Higher memory and boot cost

SYSTEM CONTAINER

OS boundary

- Shares the host kernel
- Own user space and systemd
- Fast cloning and reset

The browser labs use Incus system containers inside a controller VM.

Core operating questions transfer across distributions

Locate

- Where am I?
- Which host?
- Which user?
- Which file or unit?

Change

- Predict one state
- Use narrow privilege
- Make one edit
- Capture exit status

Verify

- Inspect resulting state
- Read relevant logs
- Test intended identity
- Prove behavior

Tools differ; the evidence loop remains stable.

Most investigations begin at five filesystem anchors

/

one root; every absolute path starts here

/etc

host-specific configuration

/home

human users' files

/var

changing data: logs, caches, queues

/usr

installed programs, libraries, shared data

Filesystem paths express intent, not physical disks

/

the root of the visible namespace

Mount

attach another filesystem at a directory

Path

a route through directory entries

File

a name referring to an inode-like object

Metadata

owner, mode, timestamps, size, and type

The hierarchy separates configuration, data, and runtime state

Configuration

- /etc host settings
- /usr shipped defaults
- Prefer drop-ins
- Track intentional edits

Persistent data

- /home user files
- /var changing service data
- /srv served content
- /opt optional software

Runtime state

- /run since boot
- /tmp temporary files
- /proc process views
- /sys device and kernel views

Directory purpose helps you search before you know an exact filename.

/proc and /sys expose live kernel state

/proc

process + kernel state

- PID directories describe processes
- cpuinfo and meminfo expose summaries
- Values may change between reads

/sys

devices + kernel objects

- Hardware and driver relationships
- Many tunables are writable by root
- Changes can affect the running system

Read freely; write only with a documented reason.

Paths are coordinates; your working directory is the origin

ABSOLUTE

```
STARTSAT /
/var/log/syslog
/etc/ssh/sshd_config
```

Works from anywhere because it begins at the root.

RELATIVE

```
STARTSAT pwd
../archive/report.txt
../scripts/check.sh
```

Meaning changes when your working directory changes.

Safety check: pwd → inspect the target → change it → verify the result

Path syntax has a small, precise vocabulary

/name	absolute: begin at root
name	relative: begin at the working directory
.	the current directory
..	the parent directory
~	shell expands to user's home
name/	directory path; trailing slash clarifies intent

Hidden names are a convention, not protection

NORMAL LISTING

ls

- Omits names beginning with dot
- Keeps routine output quieter
- Permissions are unchanged

ALL ENTRIES

ls -la

- Shows dotfiles and metadata
- Useful for .ssh and configuration
- Still governed by directory access

A dot changes default visibility; permissions control access.

The first character of ls -l identifies object type

-	regular file
d	directory
l	symbolic link
c / b	character or block device
s	Unix-domain socket
p	named pipe

Hard links and symbolic links fail in different ways

HARD LINK

another name

- Points to the same inode
- Usually cannot cross filesystems
- Data remains until last link is removed

SYMBOLIC LINK

stored path

- Can cross filesystems
- Can point to a directory
- Can become dangling

Inspect with `ls -li` and `readlink` before changing links.

Globs expand before the command starts

zero or more characters except leading dot

?

exactly one character

[ab]

one character from a set

[0-9]

one character in a range

Quoted glob

literal pattern; no pathname expansion

Inspect targets before changing files

1

Locate

pwd and realpath

2

Enumerate

printf or find the exact targets

3

Predict

state the expected names and count

4

Change

cp, mv, or rm with narrow scope

5

Verify

list, compare, or test the result

Interactive flags help, but a verified target list is stronger.

Live demo: expose a wrong-directory assumption

● ● ● PREDICT → FAIL SAFELY → CORRECT

```
whoami; hostname; pwd
ls /

ls /; cd /var/log; pwd
mkdir -p ~/demo/{in,out}
printf 'ok\nerror: disk\n' > ~/demo/in/app.log

grep "error" ~/demo/in/app.log \
| tee ~/demo/out/errors.txt
```

The error is evidence: fix the path model, not the permissions.

The shell turns text into a process request

● ● ● PROMPT + COMMAND

```
ubuntu@linux-course:~/demo$ grep -i error app.log
```

grep

program

-i

option

error

argument

app.log

argument

Exit status 0 = success · non-zero = a condition or failure

Terminal, shell, and command are separate layers

TERMINAL

input + display

- Keyboard and screen interface
- Local app or browser ttyd
- Transports control characters

SHELL

parser + launcher

- Expands variables and globs
- Connects pipelines and redirections
- Starts programs and reports status

The command is the program plus the arguments produced by the shell.

Quoting decides what the shell expands

Unquoted split words, expand variables, expand globs

'single quotes' preserve every character literally

"double quotes" expand variables but preserve one argument

backslash escape the next character

\$() replace with captured command output

Shell expansion explains surprising commands

Parse	recognize quotes, operators, and command boundaries
Parameters	replace \$name and \${name}
Commands	replace \$(...) with output
Split	unquoted results may become several words
Glob	patterns may become many pathnames
Execute	run the final program and arguments

Exit status reports success or failure

0

the command reports success

1–125

program-defined condition or failure

126

found but not executable

127

command not found

128 + signal

terminated by a signal

Control operators and pipes solve different jobs

&& AND ||

control flow

- Run next only on success
- Or run fallback on failure
- Status determines the branch

| PIPE

data flow

- Connect stdout to stdin
- Commands run as a pipeline
- Last status is reported by default

Use control operators for decisions and pipes for streams.

Navigate by asking, moving, then confirming

Where am I?

`pwd`

What is here?

`ls -la`

Move

`cd /var/log`

Create

`mkdir -p notes/archive`

Inspect

`cat · less · head · tail`

Find

`find · grep`

Use local help before trusting a random command

Quick syntax

- `command --help`
- Short option summary
- Examples may be included
- Best for recall

Full manual

- `man command`
- `/pattern` to search
- `n` for next result
- `q` to leave

Shell built-ins

- `help cd`
- `type command`
- `command -V name`
- distinguish aliases

Local documentation matches the installed version.

APT separates repositories, indexes, and installed files

Repository

- Signed metadata
- Package archives
- Release channels
- Repository trust keys

Local index

- Refreshed by apt update
- Available versions
- Candidate selection
- Dependency metadata

Installed state

- Tracked by dpkg
- Files on disk
- Package status
- APT history log

apt update refreshes knowledge; apt install changes installed state.

Use APT with an inspect-change-verify rhythm

Refresh

```
sudo apt update
```

Search

```
apt search tree
```

Inspect

```
apt show tree
```

Install

```
sudo apt install tree
```

Verify

```
dpkg -s tree; command -v tree
```

Audit

```
less /var/log/apt/history.log
```

Live demo: narrate state, not keystrokes

● ● ● BROWSER → SESSION → PIPELINE

```
whoami; hostname; pwd  
ls /
```

```
ls /; cd /var/log; pwd  
mkdir -p ~/demo/{in,out}  
printf 'ok\nerror: disk\n' > ~/demo/in/app.log
```

```
grep "error" ~/demo/in/app.log \  
| tee ~/demo/out/errors.txt
```

Predict → run → read output → verify state

Pipes connect small tools into an investigation

● ● ● ONE QUESTION, THREE FILTERS

```
journalctl -u ssh --since today | grep -i failed | tail -n 20
```

|

send output to the next command

>

replace a file with output

>>

append output to a file

Rule: inspect before ✦ — redirection happens before the command runs.

Processes start with three standard streams

stdin

file descriptor 0: keyboard or pipe input

stdout

file descriptor 1: normal output

stderr

file descriptor 2: diagnostics

>

replace a file with stdout

2>

replace a file with stderr

Build pipelines as readable evidence stages

1

Question

state the answer you
need

2

Source

choose the smallest
relevant input

3

Filter

remove unrelated
records

4

Transform

extract, sort, or count

5

Verify

inspect samples and
totals

A long pipeline is maintainable when each stage has one job and can be tested independently.

Silence does not prove success

HIDDEN ERROR

2>/dev/null

- Removes diagnostic evidence
- Can turn a real failure into silence
- Acceptable only when failure is expected

CAPTURED ERROR

2>error.log

- Preserves evidence
- Can be reviewed or attached
- Keeps stdout usable as data

Capture first; suppress only with an explicit reason.

Linux authorizes identities, not job titles

USER

UID 1001
alice

PRIMARY ♦ SUPPLEMENTARY GROUPS

alice · courseops · sudo

IDENTITY CHECKS

```
whoami
id
groups
getent passwd alice
getent group courseops
```

Processes run as a user.
Files have an owner ♦ group.
Access checks combine both.

Names are convenient; numeric IDs drive authorization

Username	human-readable account label
UID	numeric user identity used by the kernel
Group name	human-readable collaboration label
GID	numeric group identity stored on files
Effective IDs	identity currently used for access checks

Account databases separate public identity from secrets

/etc/passwd

- name and UID
- primary GID
- home directory
- login shell

/etc/group

- group name and GID
- supplementary members
- not every primary member
- readable account data

/etc/shadow

- password hashes
- aging fields
- root-readable only
- never copy casually

Use `getent` so local files and network identity sources are queried consistently.

Existing shells keep old group credentials

ACCOUNT DATABASE

updated now

- usermod changes membership
- getent shows the new record
- Future sessions inherit it

CURRENT SHELL

old credentials

- Keeps groups from login time
- id may not show the change
- Re-login or newgrp refreshes context

Always test membership in the process that will do the work.

Use sudo for one deliberate action

Normal shell

least privilege
clear ownership
smaller blast radius

✦ **sudo**

One privileged command

● ● ● EXPLICIT INTENT

```
sudo apt install tree  
sudo systemctl restart ssh  
sudoedit /etc/ssh/sshd_config
```

Avoid living in a root shell. Ask: what exact command needs privilege?

sudo policy answers who may run what as whom

Caller

the authenticated invoking user

Target

root by default or selected with -u

Command

the permitted executable and arguments

Policy

rules from sudoers and included files

Audit

logs record the privileged request

Permissions are three triplets evaluated in order

- **rwX** **r-X** - - -

OWNER

GROUP

OTHERS

rwX

r-X

- - -

file: read · change · run directory: list names · create/remove · enter/traverse

Access checks select one permission class

Owner match

use only the owner triplet

Else group match

use only the group triplet

Else

use the other triplet

No fallback

failed owner bits do not try group bits

Privilege

root and capabilities can alter normal rules

Directory r, w, and x control different abilities

Read · r

- list entry names
- does not reveal metadata alone
- ls may show names only
- needs x for useful lookup

Write · w

- create entries
- remove entries
- rename entries
- usually requires x too

Execute · x

- traverse the directory
- look up named entries
- reach deeper paths
- not 'run the directory'

For directories, x means traversal and w changes names in that directory.

Parent directories participate in every path lookup

/	traverse root
/srv	traverse service data
/srv/team	traverse team directory
file	apply file read or write bits
namei -l	display every component and its mode

4 ♦ 2 ♦ 1 makes numeric modes readable

r = 4

w = 2

x = 1

640 = rw- r-- ---

● ● ● SAME INTENT, VERIFIED

```
chmod 640 report.txt
chmod u=rw,g=r,o= report.txt
stat -c "%A %a %U:%G %n" report.txt
```

Symbolic chmod states exactly who and what changes

u	owner class
g	group class
o	other class
✦ / -	add or remove selected bits
=	set the selected class exactly
chmod g✦r	add group read without changing other bits

umask removes permissions; it never grants

REQUESTED MODE

program default

- Files commonly request 666
- Directories commonly request 777
- Execute is not added to ordinary files

UMASK

bits to clear

- 022 yields files 644
- 022 yields directories 755
- Process-local and inherited

Final mode = requested permissions with masked bits removed.

Shared work needs the right group, not world access

TEAM SHARE

```
sudo install -d \  
  -o root -g courseops -m 2770 \  
  /srv/course-share  
  
sudo chown alice:courseops status.txt  
sudo chmod 640 status.txt
```

2

setgid on a directory

New files inherit the directory group, keeping collaboration predictable.

Owner ♦ group ♦ mode must tell the same story.

setgid keeps shared directories in one group

Directory owner

usually an administrator or service account

Directory group

the collaborating team

Mode 2xxx

setgid bit on the directory

New entries

inherit the directory group

Verification

stat and create a file as a real member

Sticky protects names in shared directories

WITHOUT STICKY

shared write

- Any writer may remove entries
- File ownership does not protect names
- Easy for users to disrupt each other

WITH STICKY

mode 1xxx

- Only owner, directory owner, or root removes
- Typical example is /tmp
- Shown as t in the other execute position

Sticky controls deletion and rename, not file confidentiality.

ACLs extend permissions without replacing the base model

MODE BITS

three classes

- Owner, group, other
- Visible in `ls -l`
- Portable baseline


POSIX ACL

named entries

- Additional users or groups
- Mask limits effective named rights
- Inspect with `getfacl`

Use ACLs deliberately and document why the exception exists.

Nano is enough to make a safe first edit



```

NANO

nano ~/linux-course/notes/day1.txt

^O Write Out / save
^X Exit
^W Search
^K Cut line
^U Paste line
    
```

^ means Ctrl

- Open the exact path
- Make one small change
- Save, exit, then verify
- Use sudoedit for protected files

Editing is not verification.

Safe edits use backup, diff, and validation

1

Inspect

read the current file and
ownership

2

Protect

copy or use version
control where
appropriate

3

Edit

nano or sudoedit the
exact path

4

Validate

use the program's
syntax checker

5

Apply

reload or restart, then
verify behavior

Saving a file proves only that bytes changed. It does not prove the configuration is valid or active.

Live demo: repair access with the smallest change

● ● ● OBSERVE → EXPLAIN → REPAIR → RETEST

```
id alice
stat -c "%A %a %U:%G %n" /srv/demo-share/report.txt
namei -l /srv/demo-share/report.txt

# observe failure as another identity
# repair owner/group/mode
# repeat the exact read
```

Never begin with chmod 777.

Diagnose permission failures from identity to behavior

1

Identity

id and groups for the
failing process

2

Path

namei -l every
component

3

Object

stat owner, group,
mode, ACL

4

Policy

sudo, service user,
mount, or security layer

5

Retest

repeat as the intended
non-root identity

Changing mode before identifying the selected class destroys evidence and often widens access unnecessarily.

SSH authenticates both ends of the conversation

CLIENT

“Is this really my server?”

Checks the server host key
against `known_hosts`



SERVER

“Is this user authorized?”

Checks a signature against
`authorized_keys`

Encryption without host verification can still connect you to an attacker.

SSH separates transport, host, and user trust

Transport

encrypts and protects session integrity

Host key

proves which server endpoint answered

User proof

proves which account may log in

Account policy

server maps proof to account policy

Session

starts command, shell, forwarding, or subsystem

Host keys identify servers; user keys identify people

HOST KEY

server identity

- Private half stays on server
- Client stores public fingerprint
- Change requires investigation

USER KEY

user identity

- Private half stays with user
- Server stores public key
- Can be revoked per account

Never copy either private key to the opposite side.

Verify host keys through an independent channel

1**Obtain**

get the expected
fingerprint from an
administrator

2**Connect**

observe the presented
fingerprint

3**Compare**

match algorithm and
complete fingerprint

4**Accept**

store the verified key in
known_hosts

5**Recheck**

treat later changes as an
event

Typing yes without comparison is trust on first use, not verification.

The private key proves; the public key permits

PRIVATE KEY

`~/.ssh/id_ed25519`

- Stays on the client
- Protect with a passphrase
- Never email or copy to server

PUBLIC KEY

`~/.ssh/id_ed25519.pub`

- May be copied to servers
- Stored in `authorized_keys`
- Can be revoked independently

● ● ● GENERATE ON THE CLIENT

```
ssh-keygen -t ed25519 -a 64 -C "linux-course"
```

A passphrase and agent reduce private-key exposure

Phrase

a passphrase encrypts the private key file at rest

KDF rounds

increase cost of password guessing

Agent

holds an unlocked key for a limited session

Permissions

prevent other local users from reading the file

Rotation

replace compromised or obsolete credentials

OpenSSH checks every component of the authorization path

Home directory	owned by the account; not broadly writable
----------------	--

~/.ssh	normally mode 700
--------	-------------------

authorized_keys	normally mode 600
-----------------	-------------------

Public key line	complete algorithm and key data
-----------------	---------------------------------

Account	login shell and policy must allow access
---------	--

Server log	explains rejected ownership or mode
------------	-------------------------------------

Put safe SSH defaults in a named client profile

~/.ssh/config

```
Host linux-course
  HostName 192.168.64.10
  User trainee
  IdentityFile ~/.ssh/linux_course_ed25519
  IdentitiesOnly yes
  StrictHostKeyChecking accept-new
  ServerAliveInterval 60
```

- One alias replaces repeated flags
- IdentitiesOnly avoids key roulette
- Verify the first host fingerprint
- Use “yes” after onboarding for strictness

Same OpenSSH syntax on macOS and Windows.

SSH uses the first value obtained

Command line

explicit options for this invocation

User config

~/.ssh/config patterns in order

System config

/etc/ssh/ssh_config defaults

Host patterns

specific entries should appear before broad ones

Effective view

ssh -G alias prints resolved settings

Live demo: verify before you trust

● ● ● SERVER FINGERPRINT + USER KEY

```
mkdir -p ~/.ssh && chmod 700 ~/.ssh
ssh-keygen -t ed25519 -f ~/.ssh/id_ed25519 -N ''

ssh-keyscan -t ed25519 -p 2222 localhost > ~/host.scan
ssh-keygen -lf ~/host.scan -E sha256 # compare trusted value
install -m 600 ~/host.scan ~/.ssh/known_hosts
ssh-copy-id -i ~/.ssh/id_ed25519.pub -p 2222 operator@localhost

ssh -o BatchMode=yes -p 2222 \
  operator@localhost 'id; hostname'
```

Private key: client only. Public key: server. Host key: verify out of band.

Test SSH failures one layer at a time

1**Network**

resolve name and reach
the port

2**Host trust**

inspect `known_hosts`
and fingerprint

3**Client choice**

`ssh -G` shows user, key,
and port

4**User auth**

`ssh -vv` reveals offered
and accepted methods

5**Server policy**

read `ssh` service logs
and account state

Verbose output is most useful after the basic target, port, and identity are confirmed.

Linux logs live in files and in the systemd journal

TEXT FILES

`/var/log`

- Application-specific files
- APT and dpkg history
- Rotation changes filenames

STRUCTURED JOURNAL

`journalctl`

- Unit, boot, PID, priority fields
- Time-aware filtering
- May be volatile or persistent

Start with the service and time window—not “all logs”.

A useful log record answers five questions

When

timestamp with timezone or boot context

Where

host, container, or node

Who

unit, process, PID, or identity

What

event and severity

Why next

correlation field, exit code, or causal detail

Filter the journal toward one answer

Which service? `journalctl -u ssh`

Which time? `--since "-10 min"`

Which severity? `-p warning`

How much? `-n 50 --no-pager`

Live? `-f`

journalctl can narrow by unit, boot, time, and priority

Unit	<code>-u ssh.service</code>
Current boot	<code>-b</code>
Previous boot	<code>-b -1</code>
Time	<code>--since '-10 minutes'</code>
Priority	<code>-p warning</code>
Stable output	<code>-n 50 --no-pager -o short-iso</code>

Filter narrowly, then widen one dimension at a time

Bound

- Choose one unit
- Select one boot
- Set the time window
- Name the incident

Read

- Follow timestamps
- Find the first cause
- Capture exit reason
- Note dependencies

Widen

- Extend time
- Add priorities
- Inspect related units
- Inspect the process

Do not begin with all logs and search by intuition.

Live demo: ask the journal a smaller question

● ● ● UNIT → TIME → PRIORITY → FOLLOW

```
sudo systemctl restart ssh  
systemctl status ssh --no-pager
```

```
sudo journalctl -u ssh -n 10 --no-pager  
sudo journalctl -u ssh --since "-5 min" -p warning --no-pager  
sudo journalctl -u ssh -f
```

Read timestamps and the first causal message aloud.

History and live follow serve different jobs

BOUNDED QUERY

incident review

- Repeatable output
- Easy to save and compare
- Best for finding first cause

FOLLOW · -f

live observation

- Shows new entries
- Useful while reproducing
- Stop with Ctrl+C

Capture a bounded baseline before entering follow mode.

systemd manages declared units toward a state

UNIT FILE

What should run?

dependencies · user · command
restart policy · boot target



MANAGER

What is running?

start · stop · supervise
track exit status



JOURNAL

What happened?

events + output

“Enabled” configures future activation. “Active” describes runtime state. They are independent.

Unit files come from layered locations

Vendor units

vendor path varies by distribution

Admin units

/etc/systemd/system and overrides

Runtime units

/run/systemd/system

Drop-ins

name.service.d/*.conf overrides selected settings

Effective view

systemctl cat name.service

Service transitions leave journal evidence

Declared

- ExecStart command
- Service user
- Dependencies
- Restart policy

Transition

- Conditions
- Ordering
- Credentials
- Exit status

Observed

- Active state
- Substate
- Main PID
- Journal events

systemctl status is a current summary; journalctl provides the longer event sequence.

Requirement and ordering are separate relationships

Requires=	start together; failure relationship
Wants=	weaker activation relationship
After=	ordering only; does not pull a unit in
Before=	inverse ordering relationship
PartOf=	propagate selected lifecycle actions

Use systemctl verbs deliberately

Observe `status` · `is-active` · `is-enabled`

Change runtime `start` · `stop` · `restart` · `reload`

Change boot `enable` · `disable`

Inspect config `cat` · `show`

After unit edits `daemon-reload`

Restart a service only after capturing evidence.

Enabled, active, reloaded, and restarted are independent facts

BOOT + RUNTIME

enable / start

- enable changes boot links
- start changes current state
- --now requests both operations

CONFIG APPLY

reload / restart

- reload asks process to reread
- restart replaces runtime process
- daemon-reload rereads unit files

Name the state you intend to change before choosing the verb.

Live demo: status is the symptom; logs reveal the cause

● ● ● FAILED SERVICE → VERIFIED ENDPOINT

```
systemctl status demo-web --no-pager
sudo journalctl -u demo-web --since "-5 min" --no-pager
systemctl cat demo-web

# repair missing user + working directory
sudo systemctl restart demo-web
systemctl is-active demo-web
curl -fsS http://localhost:8090
```

State ♦ logs ♦ configuration ♦ behavior

Diagnose services from symptom to verified behavior

1**State**

systemctl status and is-active

2**Events**

journalctl by unit, boot, and time

3**Definition**

systemctl cat and show

4**Repair**

change one evidenced cause

5**Behavior**

test the endpoint as a real client

A green active state proves that systemd sees a running process, not that users receive the intended service.

Live demo: separate the first cause from later noise

Status	<code>systemctl status course-web --no-pager</code>
Events	<code>journalctl -u course-web -b --since '-10 min'</code>
Definition	<code>systemctl cat course-web</code>
Repair	change only the evidenced cause
Restart	systemctl restart course-web
Verify	is-active plus <code>curl http://127.0.0.1:8088/</code>

A safe operator closes every change with evidence

1

Locate

browser or session ·
user · directory

2

Predict

what state should
change?

3

Act

one scoped command

4

Observe

output · status · logs

5

Verify

behavior through
another check


Pair prompt: explain this loop using the failed-service demo.

HANDS-ON BLOCK

**Three hours.
One browser.
Eight incidents.**

Work in pairs when useful. Driver predicts; navigator checks the evidence.

BROWSER LAB: open your assigned seat link




01	02
03	04
05	06
07	08

Labs 1–2 establish context, then test path behavior

LAB 1 · 15 MIN

System orientation



A terminal window with a dark background and light text. The title bar shows three colored circles (red, yellow, green) followed by the text 'CONTEXT → PROFILE'. The terminal contains the following commands:

```
whoami; hostname; pwd  
. /etc/os-release; echo "$NAME"  
uname -r; uname -m  
ps -p 1 -o comm=
```

LAB 2 · 25 MIN

Filesystem ♦ links

- Build notes · archive · links
- Copy and move with verification
- Compare hard-link inode numbers
- Resolve a relative symbolic link

Checkpoint: the profile proves context; both link types behave as predicted.

Labs 3–4 turn text and packages into inspectable evidence

LAB 3 · 20 MIN

Editor ♦ text evidence

● ● ● EDIT → FILTER → COUNT

```
nano ~/app.conf  
awk '{print $4}' ~/access.log  
sort | uniq -c  
grep '^level=ERROR '
```

LAB 4 · 20 MIN

Package lifecycle

- Refresh and inspect the candidate
- Install through APT
- Map binary to owning package
- Prove behavior with tree output

Checkpoint: config, counts, package state, and command output all agree.

Labs 5–6 prove least privilege and verified remote trust

LAB 5 • 30 MIN

Permission incident

- Inspect identity + every path component
- Repair group ownership + modes
- Prove setgid inheritance with a file
- Test Alice, Bob, and outsider

LAB 6 • 25 MIN

SSH trust setup

SERVER TRUST → USER TRUST

```
chmod 700 ~/.ssh  
collect Ed25519 host key  
compare trusted SHA-256 fingerprint  
install verified known_hosts entry  
BatchMode key-only login
```

Checkpoint: team access is narrow; SSH host trust is fingerprint-verified and key-only.

Labs 7–8 move from scoped evidence to durable repair

LAB 7 • 20 MIN

Journal investigation

- Read ActiveState, Result, and status
- Scope journal by unit + boot
- Save the causal evidence
- Submit five exact facts

LAB 8 • 25 MIN

Broken service capstone

EVIDENCE → REPAIR → VERIFY

```
systemctl status course-web --no-pager  
journalctl -u course-web -b --no-pager  
install -d -o courseapp -g courseapp  
systemctl enable --now course-web  
curl -fsS http://127.0.0.1:8088/
```

Final checkpoint: cause explained; state, enablement, socket, and HTTP verified.

When blocked, explain the evidence chain

Use this five-step peer check:

- Which user, host, and working directory?
- What target state did you inspect?
- Can you repeat the exact symptom?
- Which log or exit status identifies the cause?
- What one change did you retest?

A green checker starts the explanation; it does not replace it.

Keep your seat link; keep practising the loop

NEXT 24 HOURS

- Reset Lab 8 and solve without hints
- Rebuild the SSH profile from memory
- Explain 640 and 2770 aloud

NEXT 30 DAYS

- Continue with Linux Upskill Challenge
- Read one man page per session
- Automate only after the manual path is clear

Locate → predict → act → observe → verify

Sources and open-course attribution

OPEN COURSE INSPIRATION

Linux Upskill Challenge · CC BY 4.0
linuxupskillchallenge.com

Linux Training + Linux101Labs
coverage and exercise-pattern review

PRIMARY REFERENCES

Canonical Ubuntu Server documentation
Incus system-container documentation
ttyd browser-terminal project
OpenSSH + systemd manuals

Visual direction

Platform source and delivery notes: [browser-lab-platform/README.md](#)

Full URLs, licensing notes, and adaptation details: [resources/SOURCES.md](#)